

CS136

CP3 Pengcheng Xu


Collaboration Statement:

- Total hours spent: 8 hrs
- Consult Resources:
 - Course's website
 - Numpy website
 - Online resource
- Internet

Problem 1.

1a

After implementing `calc_tart_log_pdf` and `main` block in `run_RW_prob1.py`, we get the following figure:



1b


Given the info, I'm not very confident that our MCMC chain has converged to the target distribution.

The reason is that the accept rate is a little too high (i.e. around 0.8 after 10,000 iterations). The situation is similar to that in **1a** where the $rw_stddev = 0.010$. *Even though in the current setting, the $rw_stddev = 100.0$, this value could be still too small compared to the target distribution's standard deviation (e.g. what if our target std is 100,000).* In other words, the difference (i.e. randomwalk step) between our old Sample (i.e. z_old) and new Sample (i.e. z_new) is too small, so we're sort of stuck in the only small portion of the target distribution (and get a high accept rate), but can't see the whole picture just like what we did in **figure 1a** where $rw_stddev = 0.010$.

Problem 2.

2a

The posterior predictive visualization for order-2 polynomial model:



2b

The following is the table of per-example score of models with order 0 and order 2 on the provided test set.

order	test score
0	-4.533
2	-4.220

2c

The following is the screenshot of my implementation of `calc_score` function

```
def calc_score(list_of_z_D, phi_RM, t_R):
    ''' Calculate per-example score averaged over provided test set of size R

    Args
    ----
    list_of_z_D : list of ndarray
        List of samples of parameters, assumed to be from target posterior
    phi_RM : 2D array, shape (R, M)
        Feature vectors for each of the examples in test set of size R
    t_R : 1D array, shape (R,)
        Output values for each of the examples in test set of size R

    Returns
    ----
    score : float
        Per-example log pdf of all t values in test set
        using Monte-Carlo approximation to marginal likelihood
    '''
    S = len(list_of_z_D)
    res = []
    for ss in range(S):
        z_ss_D = list_of_z_D[ss]
        # Compute score formula for ss-th sample (see instructions)
        # Hint: Use unpack_mean_N_and_stddev_N
        mean_R, stddev_R = unpack_mean_N_and_stddev_N(z_D = z_ss_D, phi_NM = phi_RM)
        res.append(norm.logpdf(t_R, mean_R, stddev_R))
    # TODO aggregate across all S samples
    # Hint: use scipy.special.logsumexp to be numerically stable
    res = np.vstack(res)
    res = scipy.special.logsumexp(res, axis = 0) - np.log(S)
    return np.average(res)
```